

Exp:1 Techniques for Data Pre-processing: Mean Removal, Scaling, Normalization

Required Libraries:

pip install pandas

pip install -U scikit-learn

CODE:

```
import pandas as pd
import numpy as np
from sklearn.preprocessing import StandardScaler

file_path = 'data.csv'
data = pd.read_csv(file_path, header=None, names=['Values'])
print("Original Data:")
print(data)

mean_values = np.mean(data['Values'])
data['Mean_Removed_Values'] = data['Values'] - mean_values

normalized_values = (data['Values'] - np.min(data['Values'])) / (np.max(data['Values']) -
np.min(data['Values']))

data['Normalized_Values'] = normalized_values

scaler = ((data['Values'] - max(data['Values']))/max(data['Values']))

print("\nMean Removed Data:")
print(data[['Mean_Removed_Values']])

print("\nNormalized Data:")
print(data[['Normalized_Values']])

print("\nStandardized Data:")
print(scaler)
```

2.a. Naïve Bayes Classifier

```
import pandas as pd

f = pd.DataFrame({'Weather':['Sunny', 'Rainy', 'Sunny', 'Sunny'],
                  'Wind':['Mild', 'Mild', 'High', 'Mild'],
                  'Temp':['Moderate', 'Mild', 'Moderate', 'Mild'],
                  'go':['Yes', 'No', 'Yes', 'Yes']})

print(f.columns)

from sklearn.naive_bayes import GaussianNB as g
from sklearn.preprocessing import LabelEncoder as le
from sklearn.model_selection import train_test_split as tt

l = le()
for i in f.columns:
    f[i] = l.fit_transform(f[i])

x = f.iloc[:, :3]
y = f.iloc[:, 3]

xtr, xte, ytr, yte = tt(x, y, test_size=0.3)

gg = g()
gg.fit(xtr, ytr)
y_pred = gg.predict(xte)

from sklearn.metrics import accuracy_score
print(accuracy_score(yte, y_pred))
```

2.b. Support vector machine

Required Libraries:

pip install matplotlib

CODE:

```
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd

dataset = pd.read_csv("Social_Network_Ads.csv")
X = dataset.iloc[:, [2, 3]].values
y = dataset.iloc[:, 4].values

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.25, random_state = 0)

from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)

from sklearn.svm import SVC
classifier = SVC(kernel='rbf', random_state = 0)
classifier.fit(X_train, y_train)
y_pred = classifier.predict(X_test)

from sklearn.metrics import confusion_matrix, accuracy_score
cm = confusion_matrix(y_test, y_pred)
print(cm)
accuracy_score(y_test, y_pred)

from matplotlib.colors import ListedColormap

X_set, y_set = X_test, y_test
X1, X2 = np.meshgrid(np.arange(start = X_set[:, 0].min() - 1, stop = X_set[:, 0].max() + 1, step =
0.01), np.arange(start = X_set[:, 1].min() - 1, stop = X_set[:, 1].max() + 1, step = 0.01))
plt.contourf(X1, X2, classifier.predict(np.array([X1.ravel(), X2.ravel()]).T).reshape(X1.shape), alpha =
0.75, cmap = ListedColormap(('red', 'green')))
plt.xlim(X1.min(), X1.max())
```

```

plt.ylim(X2.min(), X2.max())
for i, j in enumerate(np.unique(y_set)):
    plt.scatter(X_set[y_set == j, 0], X_set[y_set == j, 1], color = ListedColormap(('pink', 'green'))(i), label
= j)
plt.title('SVM (Test set)')
plt.xlabel('Age')
plt.ylabel('Estimated Salary')
plt.legend()
plt.show()

```

2.c.Logistic regression

```

import pandas as pd

import numpy as np

import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score
from matplotlib.colors import ListedColormap

dataset = pd.read_csv("diabetes.csv")

x = dataset.iloc[:, [4, 7]].values
y = dataset.iloc[:, 8].values

xtrain, xtest, ytrain, ytest = train_test_split(x, y, test_size = 0.25, random_state = 0)

sc_x = StandardScaler()

```

```

xtrain = sc_x.fit_transform(xtrain)
xtest = sc_x.transform(xtest)

print (xtrain[0:10, :])

classifier = LogisticRegression(random_state = 0)
classifier.fit(xtrain, ytrain)
y_pred = classifier.predict(xtest)

cm = confusion_matrix(ytest, y_pred)

print ("Confusion Matrix : \n", cm)

print ("Accuracy : ", accuracy_score(ytest, y_pred))

X_set, y_set = xtest, ytest
X1, X2 = np.meshgrid(np.arange(start = X_set[:, 0].min() - 1,
                             stop = X_set[:, 0].max() + 1, step = 0.01),
                     np.arange(start = X_set[:, 1].min() - 1,
                             stop = X_set[:, 1].max() + 1, step = 0.01))

plt.contourf(X1, X2, classifier.predict(
    np.array([X1.ravel(), X2.ravel()]).T).reshape(
    X1.shape), alpha = 0.75, cmap = ListedColormap(('red', 'green')))

plt.xlim(X1.min(), X1.max())
plt.ylim(X2.min(), X2.max())

for i, j in enumerate(np.unique(y_set)):

```

```
plt.scatter(X_set[y_set == j, 0], X_set[y_set == j, 1],  
            c = ListedColormap(('red', 'green'))(i), label = j)
```

```
plt.title('Classifier (Test set)')
```

```
plt.xlabel('Age')
```

```
plt.ylabel('Diabetes')
```

```
plt.legend()
```

```
plt.show()
```

2.d.DECISION TREE

```
import pandas as pd
```

```
import math
```

```
import numpy as np
```

```
data = pd.read_csv("3-dataset.csv")
```

```
features = [feat for feat in data]
```

```
features.remove("answer")
```

```
class Node:
```

```
    def __init__(self):
```

```
        self.children = []
```

```
        self.value = ""
```

```
        self.isLeaf = False
```

```
        self.pred = ""
```

```
    def entropy(examples):
```

```
        pos = 0.0
```

```
        neg = 0.0
```

```
        for _, row in examples.iterrows():
```

```
            if row["answer"] == "yes":
```

```

        pos += 1
    else:
        neg += 1
    if pos == 0.0 or neg == 0.0:
        return 0.0
    else:
        p = pos / (pos + neg)
        n = neg / (pos + neg)
        return -(p * math.log(p, 2) + n * math.log(n, 2))

```

```

def info_gain(examples, attr):
    uniq = np.unique(examples[attr])
    gain = entropy(examples)
    for u in uniq:
        subdata = examples[examples[attr] == u]
        sub_e = entropy(subdata)
        gain -= (float(len(subdata)) / float(len(examples))) * sub_e
    return gain

```

```

def ID3(examples, attrs):
    root = Node()

    max_gain = 0
    max_feat = ""
    for feature in attrs:
        gain = info_gain(examples, feature)
        if gain > max_gain:
            max_gain = gain
            max_feat = feature
    root.value = max_feat

```

```

uniq = np.unique(examples[max_feat])
for u in uniq:
    subdata = examples[examples[max_feat] == u]
    if entropy(subdata) == 0.0:
        newNode = Node()
        newNode.isLeaf = True
        newNode.value = u
        newNode.pred = np.unique(subdata["answer"])
        root.children.append(newNode)
    else:
        dummyNode = Node()
        dummyNode.value = u
        new_attrs = attrs.copy()
        new_attrs.remove(max_feat)
        child = ID3(subdata, new_attrs)
        dummyNode.children.append(child)
        root.children.append(dummyNode)

return root

```

```

def printTree(root: Node, depth=0):

```

```

    for i in range(depth):
        print("\t", end="")
    print(root.value, end="")
    if root.isLeaf:
        print(" -> ", root.pred)
    print()
    for child in root.children:
        printTree(child, depth + 1)

```

```

def classify(root: Node, new):

```



```

for child in root.children:
    if child.value == new[root.value]:
        if child.isLeaf:
            print ("Predicted Label for new example", new, " is:", child.pred)
            exit
        else:
            classify (child.children[0], new)

```

```

root = ID3(data, features)
print("Decision Tree is:")
printTree(root)
print ("-----")

```

```

new = {"outlook": "sunny", "temperature": "hot", "humidity": "normal", "wind": "strong"}
classify (root, new)

```

```

===== RESTART: C:\Users\online\Desktop\ml\12 decision\decision tree.py =====
Decision Tree is:
outlook
  overcast -> ['yes']
  rain
    wind
      strong -> ['no']
      weak -> ['yes']
  sunny
    humidity
      high -> ['no']
      normal -> ['yes']
-----
Predicted Label for new example {'outlook': 'sunny', 'temperature': 'hot', 'humidity': 'normal', 'wind': 'strong'} is: ['yes']

```

2.e. Random forest

Required Libraries:

pip install seaborn

CODE:

```
import pandas as pd

data=pd.read_csv("heart.csv")

X =data.iloc[:,[1,2,3,4,5,6,7,8,9,10,11,12]].values

y =data.iloc[:,13].values

from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=1)

from sklearn.ensemble import RandomForestClassifier

rfc=RandomForestClassifier()

rfc.fit(X_train, y_train)

y_pred=rfc.predict(X_test)

from sklearn import metrics

print("Classification Accuracy:", metrics.accuracy_score(y_test, y_pred)*100)

cm=metrics.confusion_matrix(y_test,y_pred)

print(cm)

import seaborn as sn

from matplotlib import pyplot as plt

plt.figure(figsize=(5,4))

sn.heatmap(cm,annot=True)

plt.xlabel('Predicted value')

plt.ylabel('Actual value')

plt.show()
```

3.K-MEANS

```
import pandas as pd

import numpy as np

import matplotlib.pyplot as plt

import seaborn as sns

from sklearn.cluster import KMeans

from sklearn.metrics import silhouette_score

from sklearn.preprocessing import MinMaxScaler

iris = pd.read_csv("Iris1.csv")

x = iris.iloc[:, [1,2,3,4]].values


from sklearn.cluster import KMeans

wcss = []


for i in range(1, 11):

    kmeans = KMeans(n_clusters = i, init = 'k-means++', max_iter = 300, n_init = 10, random_state = 0)

    kmeans.fit(x)

    wcss.append(kmeans.inertia_)


kmeans = KMeans(n_clusters = 3, init = 'k-means++', max_iter = 300, n_init = 10, random_state = 0)

y_kmeans = kmeans.fit_predict(x)

plt.scatter(x[y_kmeans == 0, 0], x[y_kmeans == 0, 1], s = 100, c = 'blue', label = 'Iris-setosa')

plt.scatter(x[y_kmeans == 1, 0], x[y_kmeans == 1, 1], s = 100, c = 'orange', label = 'Iris-versicolour')

plt.scatter(x[y_kmeans == 2, 0], x[y_kmeans == 2, 1], s = 100, c = 'green', label = 'Iris-virginica')

plt.scatter(kmeans.cluster_centers_[0, 0], kmeans.cluster_centers_[0,1], s = 100, c = 'red', label =
'Centroids')

plt.legend()

plt.show()
```

EXPERIMENT 4 NLTK

Required Libraries:

pip install nltk

CODE:

```
import nltk

from nltk.tokenize import word_tokenize

from nltk.stem import PorterStemmer, WordNetLemmatizer

from nltk.corpus import stopwords

# Download NLTK resources if not already installed
nltk.download('punkt')
nltk.download('stopwords')
nltk.download('wordnet')

# Sample text
text = "The quick brown foxes are jumping over the lazy dogs. The dogs are not amused."

# Tokenization
tokens = word_tokenize(text)

# Remove stop words
stop_words = set(stopwords.words('english'))
filtered_tokens = [word for word in tokens if word.lower() not in stop_words]

# Stemming
stemmer = PorterStemmer()
stemmed_tokens = [stemmer.stem(word) for word in filtered_tokens]

# Lemmatization
```

```

lemmatizer = WordNetLemmatizer()
lemmatized_tokens = [lemmatizer.lemmatize(word) for word in filtered_tokens]

# Display the results
print("Original Text:", text)
print("\nTokenization:", tokens)
print("\nFiltered Tokens (without stop words):", filtered_tokens)
print("\nStemmed Tokens:", stemmed_tokens)
print("\nLemmatized Tokens:", lemmatized_tokens)

```

EXPERIMENT – 5(BUILDING BAG OF WORDS MODEL USING NLTK)

```

import nltk

from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize
from nltk.probability import FreqDist

nltk.download('punkt')
nltk.download('stopwords')

def preprocess_text(text):
    stop_words = set(stopwords.words('english'))

    word_tokens = word_tokenize(text)

    filtered_words = [word.lower() for word in word_tokens if word.isalpha() and word.lower() not in
stop_words]

    return filtered_words

def create_bow_model(texts):
    all_words = []

```

```
for text in texts:
```

```
    words = preprocess_text(text)
```

```
    all_words.extend(words)
```

```
word_freq = FreqDist(all_words)
```

```
bow_model = {word: freq for word, freq in word_freq.items()}
```

```
return bow_model
```

```
# Example usage
```

```
texts = [
```

```
    "The cat sat on the mat, and the mat was comfortable.",
```

```
    "She sang a sweet song, a song that touched everyone's heart.",
```

```
    "Coding coding can be challenging, but coding is also incredibly rewarding.",
```

```
]
```

```
bow_model = create_bow_model(texts)
```

```
# Print the Bag of Words model
```

```
print("Bag of Words Model:")
```

```
for word, freq in bow_model.items():
```

```
    print(f"{word}: {freq}")
```

Exp 6: Topic Modeling: Identifying Patterns in Text Data

```
import csv
import re

def identify_patterns(csv_file_path, column_name):
    patterns = {}

    with open(csv_file_path, 'r') as csvfile:
        reader = csv.DictReader(csvfile)

        for row in reader:
            text = row[column_name]

            # Example pattern: finding words that start with 'pattern'
            pattern_matches = re.findall(r'Female', text, flags=re.IGNORECASE)

            # Update patterns dictionary with matches
            for match in pattern_matches:
                if match in patterns:
                    patterns[match] += 1
                else:
                    patterns[match] = 1

    return patterns

csv_file_path = '2b Social_Network_Ads.csv' # Update with your CSV file path
column_name = 'Gender' # Update with the actual column name in your CSV file

result = identify_patterns(csv_file_path, column_name)

# Display the identified patterns and their counts
for pattern, count in result.items():
    print(f"Pattern: {pattern}, Count: {count}")
```

Output:

Pattern: Female, Count: 204

EXPERIMENT – 7 HIDDEN MARKOV MODEL

Required Libraries:

pip install hmmlearn

CODE:

```
import numpy as np

from hmmlearn import hmm

# Step 1: Define Model Parameters

n_states = 2 # Number of hidden states (Rainy and Sunny)

# Transition matrix (A): Probability of transitioning from one state to another

trans_matrix = np.array([[0.7, 0.3], [0.4, 0.6]])

# Ensure rows of the transition matrix sum to 1

trans_matrix /= trans_matrix.sum(axis=1, keepdims=True)

# Emission matrix (B): Probability of observing an emission given the current state

emission_matrix = np.array([[0.1, 0.4, 0.5], [0.6, 0.3, 0.1]])

# Initial state probabilities ( $\pi$ ): Probability distribution of starting in each state

initial_probs = np.array([0.6, 0.4])

# Step 2: Create HMM Model

model = hmm.MultinomialHMM(n_components=n_states,

                           startprob_prior=initial_probs,

                           transmat_prior=trans_matrix,

                           n_iter=100)
```



```
# Step 3: Generate Training Data (for simplicity, you can use a pre-existing dataset)
```

```
# Observations: 0 - Umbrella, 1 - Jacket, 2 - T-shirt
```

```
train_data = np.array([[0, 1, 2, 0, 1, 2, 0, 2, 1]])
```

```
# Reshape the array if needed
```

```
train_data = train_data.reshape(-1, 1)
```

```
# Step 4: Fit the Model
```

```
model.fit(train_data)
```

```
# Step 5: Predict States for a New Sequence
```

```
new_data = np.array([[0, 2, 1]]) # Umbrella, T-shirt, Jacket
```

```
new_data = new_data.reshape(-1, 1)
```

```
predicted_states = model.predict(new_data)
```

```
# Map numerical predictions to weather states
```

```
weather_states = ['Rainy', 'Sunny']
```

```
predicted_states_text = [weather_states[state] for state in predicted_states]
```

```
# Display Results
```

```
print("Predicted Weather States:", predicted_states_text)
```

Output:

```
Predicted Weather States: ['Rainy', 'Sunny', 'Rainy']
```

EXPERIMENT - 9 A Bot to Play Tic Tac Toe

```
def print_board(board):
    for row in board:
        print(" | ".join(row))
        print("-" * 5)

def check_winner(board, player):
    # Check rows, columns, and diagonals
    for i in range(3):
        if all(board[i][j] == player for j in range(3)) or all(board[j][i] == player for j in range(3)):
            return True

    if all(board[i][i] == player for i in range(3)) or all(board[i][2 - i] == player for i in range(3)):
        return True

    return False

def is_board_full(board):
    return all(board[i][j] != " " for i in range(3) for j in range(3))

def tic_tac_toe():
    board = [[" " for _ in range(3)] for _ in range(3)]
    players = ["X", "O"]
    current_player = players[0]

    while True:
        print_board(board)

        # Get player move
        while True:
            row = int(input("Enter row (0, 1, or 2): "))
            col = int(input("Enter column (0, 1, or 2): "))

            if 0 <= row < 3 and 0 <= col < 3 and board[row][col] == " ":
                break
```

```

    else:
        print("Invalid move. Try again.")
    # Make the move
    board[row][col] = current_player
    # Check for a winner
    if check_winner(board, current_player):
        print_board(board)
        print(f"Player {current_player} wins!")
        break
    # Check for a tie
    if is_board_full(board):
        print_board(board)
        print("It's a tie!")
        break

    # Switch to the other player
    current_player = players[1] if current_player == players[0] else players[0]

if __name__ == "__main__":
    tic_tac_toe()

```

EXPERIMENT-10 SINGLE AND MULTIPLE LAYER PERCEPTRON

Required Libraries:

pip install tensorflow

CODE:

```
import numpy as np

import tensorflow as tf

from tensorflow.keras import layers, models

from sklearn.datasets import load_iris

from sklearn.model_selection import train_test_split

from sklearn.preprocessing import StandardScaler

from sklearn.metrics import accuracy_score


# Load and preprocess the Iris dataset

X, y = load_iris(return_X_y=True)

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

X_train, X_test = StandardScaler().fit_transform(X_train),
StandardScaler().fit_transform(X_test)


# Define and compile a single-layer neural network

model_single_layer = models.Sequential([layers.Dense(64, 'relu', input_shape=(4,)),
layers.Dense(3, 'softmax')])

model_single_layer.compile(optimizer='adam', loss='sparse_categorical_crossentropy',
metrics=['accuracy'])

model_single_layer.fit(X_train, y_train, epochs=15, validation_data=(X_test, y_test))


# Evaluate the single-layer model
```

```

single_layer_accuracy = accuracy_score(y_test,
np.argmax(model_single_layer.predict(X_test), axis=1))

print(f"\nSingle-layer Neural Network - Accuracy: {single_layer_accuracy}")

# Define and compile a multi-layer neural network

model_multi_layer = models.Sequential([layers.Dense(64, 'relu', input_shape=(4,)),
layers.Dense(32, 'relu'), layers.Dense(3, 'softmax')])

model_multi_layer.compile(optimizer='adam', loss='sparse_categorical_crossentropy',
metrics=['accuracy'])

model_multi_layer.fit(X_train, y_train, epochs=15, validation_data=(X_test, y_test))

# Evaluate the multi-layer model

multi_layer_accuracy = accuracy_score(y_test,
np.argmax(model_multi_layer.predict(X_test), axis=1))

print(f"\nMulti-layer Neural Network - Accuracy: {multi_layer_accuracy}")

```

EXPERIMENT – 11 BUILDING LINEAR REGRESSION USING ANN

```

import tensorflow as tf

import numpy as np

import matplotlib.pyplot as plt

# Generate some random data for demonstration

np.random.seed(0)

X_train = np.random.rand(100, 1)

y_train = 2 * X_train + 1 + 0.1 * np.random.randn(100, 1)

# Build the model

model = tf.keras.Sequential([

    tf.keras.layers.Dense(units=1, input_shape=(1,))

])

```

```
# Compile the model
model.compile(optimizer='sgd', loss='mean_squared_error')

# Train the model
history = model.fit(X_train, y_train, epochs=100, verbose=0)

# Plot the training loss over epochs
plt.plot(history.history['loss'])
plt.xlabel('Epochs')
plt.ylabel('Mean Squared Error Loss')
plt.title('Training Loss')
plt.show()

# Make predictions on new data
X_test = np.array([[0.2], [0.5], [0.8]])
predictions = model.predict(X_test)

# Print the predictions
for i in range(len(X_test)):
    print(f"Input: {X_test[i][0]}, Predicted Output: {predictions[i][0]}")
```

12. IMAGE CLASSIFIER AN APPLICATION OF DEEP LEARNING

```
import tensorflow as tf

from tensorflow.keras.models import Sequential

from tensorflow.keras.layers import Conv2D, Dense, Flatten, MaxPooling2D

# Load and preprocess the MNIST dataset

mnist = tf.keras.datasets.mnist

(X_train, y_train), (X_test, y_test) = mnist.load_data()

X_train, X_test = X_train / 255.0, X_test / 255.0 # Normalize pixel values

# Define the CNN model

model = Sequential([

    Conv2D(32, kernel_size=(3, 3), activation='relu', input_shape=(28, 28, 1)),

    MaxPooling2D(pool_size=(2, 2)),

    Conv2D(64, (3, 3), activation='relu'),

    MaxPooling2D(pool_size=(2, 2)),

    Flatten(),

    Dense(64, activation='relu'),

    Dense(10, activation='softmax')

])

# Compile the model

model.compile(optimizer='adam',

              loss='sparse_categorical_crossentropy',

              metrics=['accuracy'])

# Train the model

model.fit(X_train, y_train, epochs=5, validation_split=0.1)

# Evaluate the model

test_loss, test_acc = model.evaluate(X_test, y_test)

print(f"\nTest accuracy: {test_acc}')
```